

 New Book: [Build Your Own Database](#)

## 09. Data Serialization

For now, our server protocol response is an error code plus a string. What if we need to return more complicated data? For example, we could add the `keys` command, which returns a list of strings. We have already encoded the list-of-strings data in the request protocol. In this chapter, we will generalize the encoding to handle different types of data. This is often called “serialization”.

### 9.1 The Command Interface

Our serialization protocol consists of five data types:

```
enum {  
    SER_NIL = 0,      // Like `NULL`  
    SER_ERR = 1,      // An error code and message  
    SER_STR = 2,       // A string  
    SER_INT = 3,       // A int64  
    SER_ARR = 4,       // Array  
};
```

The array type can contain any type of data, even nested arrays.

The code listing starts with the `try_one_request` function:

```
static bool try_one_request(Conn *conn) {  
    // code omitted...  
  
    // parse the request  
    std::vector<std::string> cmd;  
    if (0 != parse_req(&conn->rbuf[4], len, cmd)) {
```

```

        msg("bad req");
        conn->state = STATE_END;
        return false;
    }

    // got one request, generate the response.
    std::string out;
    do_request(cmd, out);

    // pack the response into the buffer
    if (4 + out.size() > k_max_msg) {
        out.clear();
        out_err(out, ERR_2BIG, "response is too big");
    }
    uint32_t wlen = (uint32_t)out.size();
    memcpy(&conn->wbuf[0], &wlen, 4);
    memcpy(&conn->wbuf[4], out.data(), out.size());
    conn->wbuf_size = 4 + wlen;

    // code omitted...
}

```

For convenience, `std::string` is used to hold the response data. Production-grade projects often have more sophisticated ways to manage buffers.

The new command `keys` is added to the `do_request` handler:

```

static void do_request(std::vector<std::string> &cmd, std::string &out) {
    if (cmd.size() == 1 && cmd_is(cmd[0], "keys")) {
        do_keys(cmd, out);
    } else if (cmd.size() == 2 && cmd_is(cmd[0], "get")) {
        do_get(cmd, out);
    } else if (cmd.size() == 3 && cmd_is(cmd[0], "set")) {
        do_set(cmd, out);
    } else if (cmd.size() == 2 && cmd_is(cmd[0], "del")) {
        do_del(cmd, out);
    } else {
        // cmd is not recognized
        out_err(out, ERR_UNKNOWN, "Unknown cmd");
    }
}

```

## 9.2 Data Encoding Scheme

The code for our serialization protocol:

```
static void out_nil(std::string &out) {
    out.push_back(SER_NIL);
}

static void out_str(std::string &out, const std::string &val) {
    out.push_back(SER_STR);
    uint32_t len = (uint32_t)val.size();
    out.append((char *)&len, 4);
    out.append(val);
}

static void out_int(std::string &out, int64_t val) {
    out.push_back(SER_INT);
    out.append((char *)&val, 8);
}

static void out_err(std::string &out, int32_t code, const std::string &msg)
{
    out.push_back(SER_ERR);
    out.append((char *)&code, 4);
    uint32_t len = (uint32_t)msg.size();
    out.append((char *)&len, 4);
    out.append(msg);
}

static void out_arr(std::string &out, uint32_t n) {
    out.push_back(SER_ARR);
    out.append((char *)&n, 4);
}
```

As we can see, our serialization protocol starts with a byte of data type, followed by various types of payload data. Arrays come first with their size, then their possibly nested elements.

The serialization scheme can be summarized as “type-length-value” (TLV): “Type” indicates the type of the value; “Length” is for variable length data such as strings or arrays; “Value” is the encoded at last.

TLV is the basis of many real-world serialization protocols. It has many advantages:

1. It can be decoded without a schema, like JSON or XML, which enables some types of middleware.
2. It can encode arbitrarily nested data.

The Thrift RPC framework includes 2 serialization schemes, both derived from the TLV scheme. You can learn more by reading the specification and comparing it to the popular Protobuf scheme.

## 9.3 Command Responses

The `do_keys` function generates a response consisting of a list of strings:

```
static void h_scan(HTab *tab, void (*f)(HNode *, void *), void *arg) {
    if (tab->size == 0) {
        return;
    }
    for (size_t i = 0; i < tab->mask + 1; ++i) {
        HNode *node = tab->tab[i];
        while (node) {
            f(node, arg);
            node = node->next;
        }
    }
}

static void cb_scan(HNode *node, void *arg) {
    std::string &out = *(std::string *)arg;
    out_str(out, container_of(node, Entry, node)->key);
}

static void do_keys(std::vector<std::string> &cmd, std::string &out) {
    (void)cmd;
    out_arr(out, (uint32_t)hm_size(&g_data.db));
    h_scan(&g_data.db.ht1, &cb_scan, &out);
    h_scan(&g_data.db.ht2, &cb_scan, &out);
}
```

The `del` command responds with an integer indicating whether the deletion took place.

```
static void do_del(std::vector<std::string> &cmd, std::string &out) {
    Entry key;
```

```

key.key.swap(cmd[1]);
key.node.hcode = str_hash((uint8_t *)key.key.data(), key.key.size());

HNode *node = hm_pop(&g_data.db, &key.node, &entry_eq);
if (node) {
    delete container_of(node, Entry, node);
}
return out_int(out, node ? 1 : 0);
}

```

The code for the other commands is of nothing interesting, so we'll skip them.

## 9.4 The Client and Testing

Listing the client “deserialization” code:

```

static int32_t on_response(const uint8_t *data, size_t size) {
    if (size < 1) {
        msg("bad response");
        return -1;
    }

    switch (data[0]) {
    case SER_NIL:
        printf("(nil)\n");
        return 1;
    case SER_ERR:
        if (size < 1 + 8) {
            msg("bad response");
            return -1;
        }
        {
            int32_t code = 0;
            uint32_t len = 0;
            memcpy(&code, &data[1], 4);
            memcpy(&len, &data[1 + 4], 4);
            if (size < 1 + 8 + len) {
                msg("bad response");
                return -1;
            }
            printf("(err) %d %.*s\n", code, len, &data[1 + 8]);
            return 1 + 8 + len;
        }
    }
}

```

```

case SER_STR:
    // code omitted...

case SER_INT:
    // code omitted...

case SER_ARR:
    if (size < 1 + 4) {
        msg("bad response");
        return -1;
    }
    {
        uint32_t len = 0;
        memcpy(&len, &data[1], 4);
        printf("(arr) len=%u\n", len);
        size_t arr_bytes = 1 + 4;
        for (uint32_t i = 0; i < len; ++i) {
            int32_t rv = on_response(&data[arr_bytes], size -
arr_bytes);
            if (rv < 0) {
                return rv;
            }
            arr_bytes += (size_t)rv;
        }
        printf("(arr) end\n");
        return (int32_t)arr_bytes;
    }

default:
    msg("bad response");
    return -1;
}
}

```

Testing our new server/client:

```

$ ./client asdf
(err) 1 Unknown cmd
$ ./client get asdf
(nil)
$ ./client set k v
(nil)
$ ./client get k
(str) v
$ ./client keys

```

```
(arr) len=1
(str) k
(arr) end
$ ./client del k
(int) 1
$ ./client del k
(int) 0
$ ./client keys
(arr) len=0
(arr) end
```

Source code:

- [09\\_client.cpp](#)
- [09\\_server.cpp](#)
- [hashtable.cpp](#)
- [hashtable.h](#)

See also:

**[codecrafters.io](#)** offers “Build Your Own X” courses in many programming languages. Including Redis, Git, SQLite, Docker, and more.

---

Check it out

[← prev](#)

[Contents](#)

[next →](#)



Subscribe



Twitter



About



Free Books



Blogs



Home

Copyright © 2023 James Smith <[js@build-your-own.org](mailto:js@build-your-own.org)>